


# Performance Characteristics of the Interplanetary Overlay Network in 10 Gbps Networks

John Huff

# Hardware



Dual socket Intel Xeon CPU E5-2697 v2 @ 3GHz

24 cores (48 virtual CPUs with hyperthreading)

Network interfaces: Intel Ethernet 10-Gigabit X540-AT2

iPerf 9.41 Gb/s

# Testing methods

## TCPCLA:

- Modified version of bpdriver/bpcounter
- Average instantaneous goodput on receiver over second half of test

## CFDP:

- Time transfer of 1 GB file read from and written to ramdisk (tmpfs)


# Testing Parameters

- Convergence layer adapters: TCPCLA
- Bundle size
- TCP buffer size
- Thread CPU affinity
- Memory vs File ZCO
- CFDP

**NOTE:** Unless otherwise noted, all tests were conducted running ION on a ramdisk, max heap is at least as large as bundle size, all threads locked to same CPU, rate limiting is turned off.

# Bundle Size

ION 3.7.0 stock



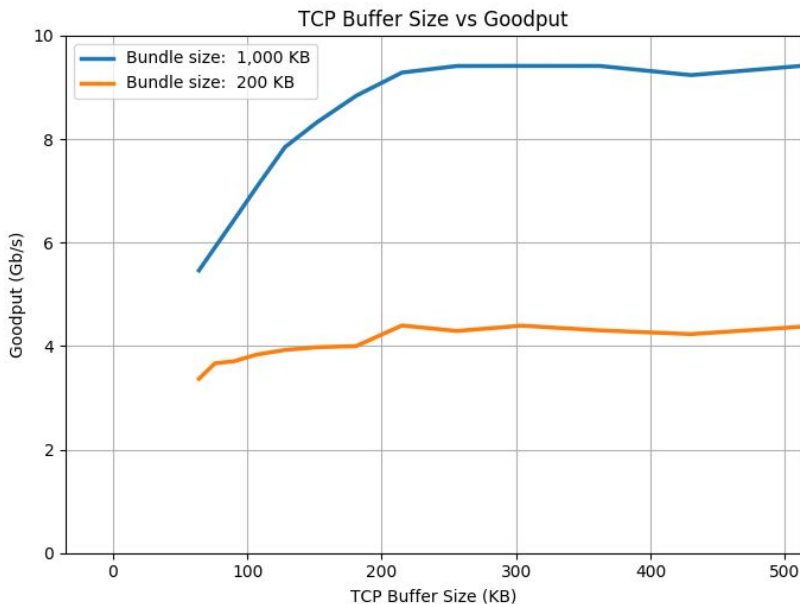
# TCP buffer size

```
bp/tcp/tcpcli.c
```


```
#define TCPCL_BUFSZ (64 * 1024)
```

Buffer to read from socket into.

Data from buffer is stored in extent and appended to the bundle ZCO. Bundle is a ZCO consisting of a list of extents, each with max length of TCP buffer size.

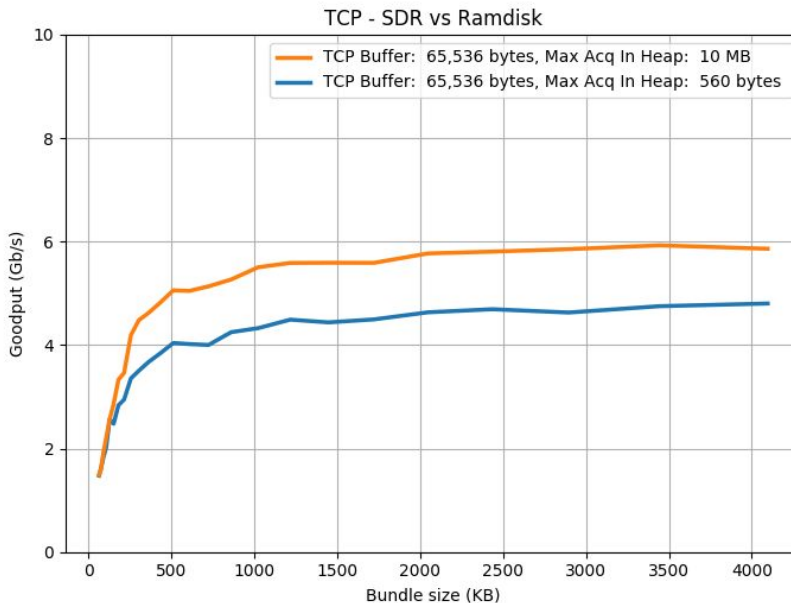


# TCP buffer size



# SDR vs Ramdisk

Using filesystem storage for ZCO decreased performance, even when using ramdisk.

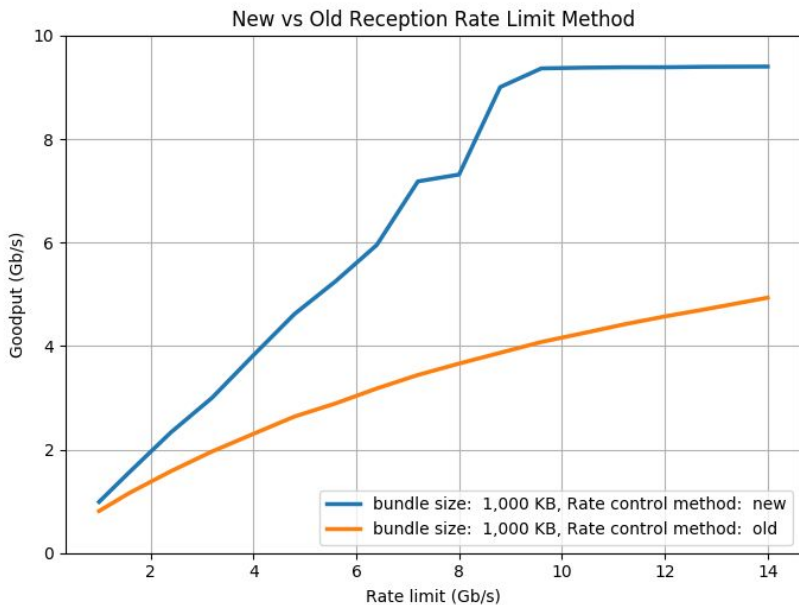




# Rate limiting

Bug causing asymptotic behavior

Using xmit rate = 0 leaves rate limiting to TCP, but has undefined behavior in contact graph routing




# CPU Affinity

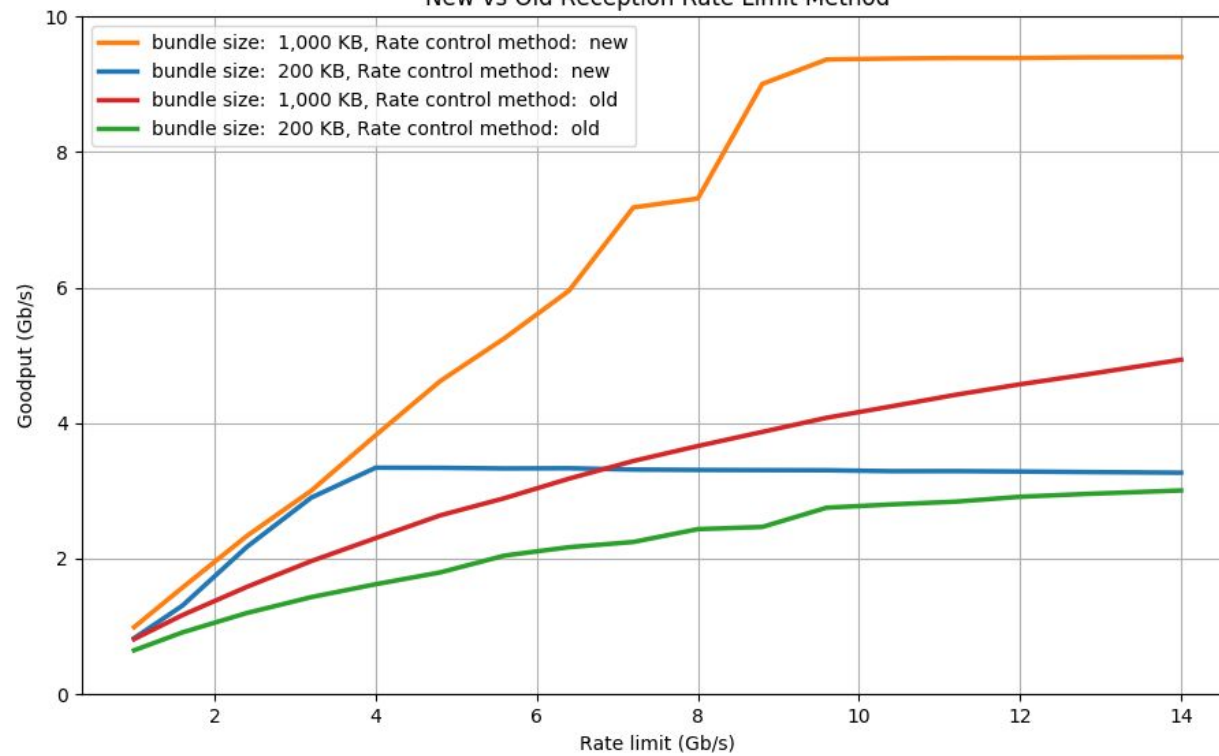
Locking all threads on all processes of an ION node greatly increases performance. (UNIX util: taskset)

Performance bump likely due to CPU caching. Same core = L1, L2 cache

Hardware dependent. Performance on hardware without caching would likely not be affected or be negatively affected.



New vs Old Reception Rate Limit Method



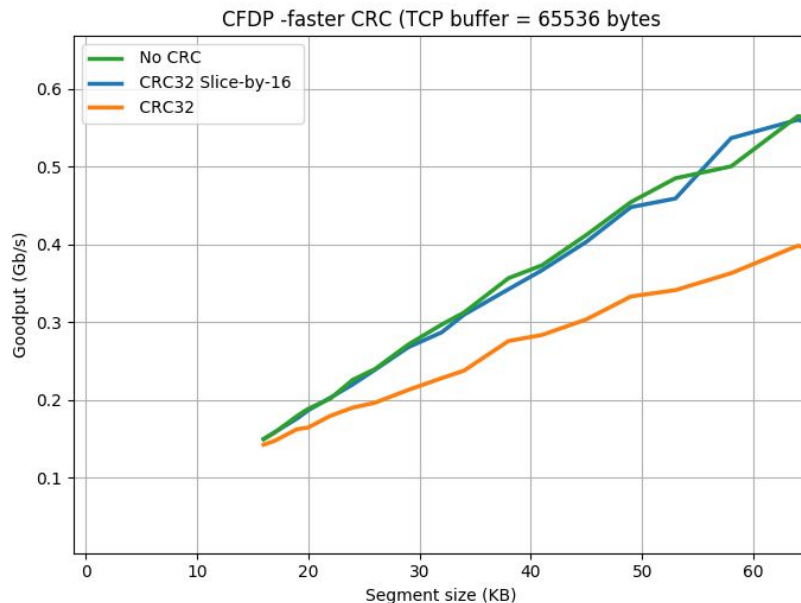
# CFDP

## CCSDS File Delivery Protocol

Used in BPCP Utility

Cyclic redundancy check (CRC)  
error-detecting code. Checksum for  
file.


Always calculated on sender, even if  
receiver doesn't request it



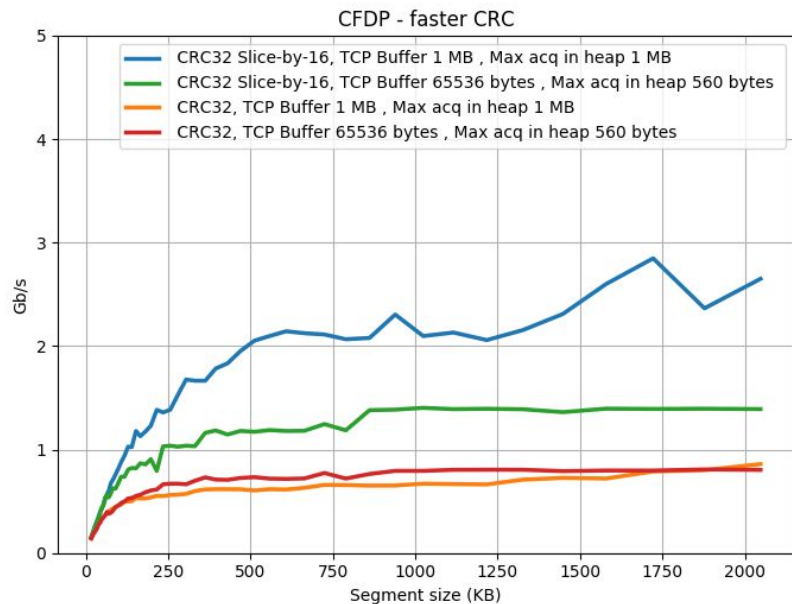
# CFDP

Modified protocol to allow for segment sizes  $> 65536$

Same result might be achieved using Delay-Tolerant Payload Conditioning (DTPC). New UT-layer adapter required.



# CFDP with larger TCP buffer



# Summary & 10Gbps Deployment Recommendations

- Use bundle sizes  $\geq 1\text{M}$  (or  $\geq 512\text{ KB}$  if using large TCP buffer)
- Use TCP buffer size  $\geq 512\text{ KB}$  (bpv\*/tcp/tcpcli.c)
- Set heapmax  $\geq$  maximum expected bundle size and  $\geq$  TCP buffer size (bpadm: m heapmax <max database heap for any single acquisition>)
- Lock all threads from all processes of same ION node to the same core. (taskset Unix util)
- CFDP: Use CRC32 slice-by-16

Repo branch: ion-4.0.0-high-speed. To compile use ./configure --enable-high-speed. Increases TCP buffer size to 512 KB, sets default CFDP checksum type to CRC slice-by-16. Addresses TCP rate limiting bug.

# Future work

Test Licklider Transmission Protocol (LTP) over UDP, TCP

Test on slower hardware (Rockchip RK3399)